

## **Project -2**

### **Data Analysis**

**By – Mitushi Vishwakarma**

**email id : [mitushi.vish@gmail.com](mailto:mitushi.vish@gmail.com)**

**Data Engineering Batch**

# ***Analysis of Global Land Average Temperature Trends***

## **Project Overview :**

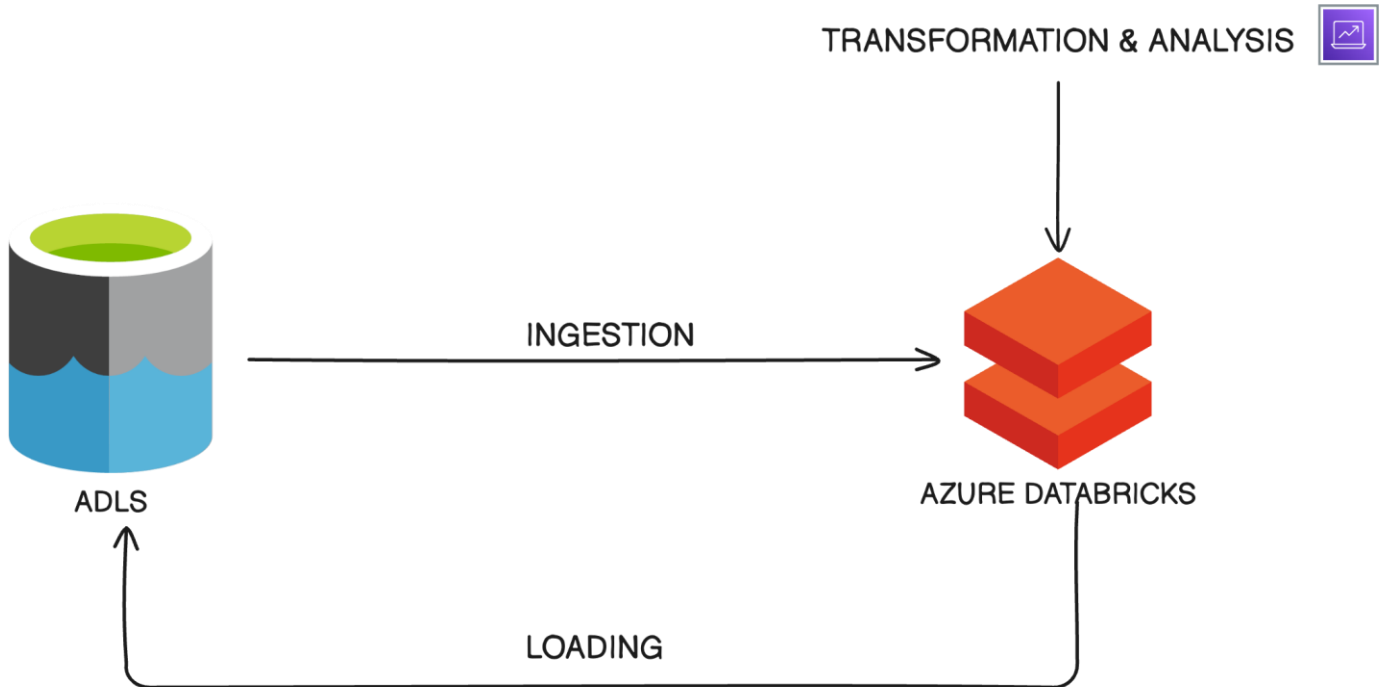
### **Analysis of Global Land Average Temperature Trends**

Our project is to implement data analysis using Spark SQL on Azure Databricks and process data for errors, seasonality, and anomalies. Some say climate change is the biggest threat of our age while others say it's a myth based on dodgy science. The goal of this project is to analyze the Global Land Average Temperature dataset to understand long-term temperature trends, detect anomalies, and identify potential factors contributing to climate change.

## **Project Requirements:**

1. **Data Source:** We need access to data that needs to be analysed. This could be stored in various formats such as CSV, Parquet, JSON, or in a database.
2. **Azure Databricks Environment:** Set up an Azure Databricks workspace and create a cluster with appropriate configurations based on the size of your data and computational requirements.
3. **PySparkSQL :** PySparkSQL, which is the Python API for Spark SQL. This includes understanding how to create SparkSession, loading data, performing data transformations, and executing SQL queries.
4. **Data Preparation:** Data often requires preprocessing before analysis. This may involve handling missing values, converting data types, filtering out irrelevant data, and ensuring the data is properly formatted for analysis.
5. **Data Analysis:**
  - Utilize PySparkSQL queries to analyze the temperature data for long-term trends, seasonal variations, and anomalies.
  - Aggregate the data to calculate statistical metrics such as average temperature by year or decade.
6. **Visualization:** Visualize the results of the analysis using plots, charts, and maps to convey insights effectively.

## Architecture :



## Azure Resources Used for this Project:

### AZURE TOOLS :

- **Azure Data Lake Storage** : Azure Data Lake Storage (ADLS) is a scalable and secure cloud-based storage solution provided by Microsoft Azure. It's designed for big data analytics workloads and is optimized for storing large amounts of structured, semi-structured, and unstructured data.
- **Azure Databricks** : Azure Databricks is a fast, easy, and collaborative Apache Spark-based analytics platform optimized for Azure. It provides a fully managed, cloud-based environment that integrates seamlessly with other Azure services, allowing data engineers, data scientists, and analysts to collaborate on big data and machine learning projects.

## **AZURE TECHNOLOGIES :**

- **Pyspark :** PySpark is the Python API for Apache Spark, a distributed computing framework for processing large datasets. It allows developers to write Spark applications using Python programming language, leveraging Spark's distributed computing capabilities.
- **Spark SQL:** Spark SQL is a module in Apache Spark for processing structured data using SQL and DataFrame API. It provides a unified interface for querying structured data sources, enabling seamless integration of SQL queries with Spark's distributed processing engine.

## **How it Works :**

1. Data Ingestion:  
Load the dataset into Azure Databricks storage (e.g., Azure Blob Storage).
2. Data Exploration and Cleaning:
  - Explore the structure and contents of the dataset to understand its schema and characteristics.
  - Perform data cleaning steps to handle missing values, outliers, and inconsistencies.
3. Data Analysis:
  - Utilize PySparkSQL queries to analyze the data for errors, seasonality, and anomalies.
  - Aggregate the data to calculate statistics such as average.
4. Error Detection:  
Utilize Spark SQL queries or built-in functions to detect errors in your data. This could involve identifying inconsistencies, duplicates, or unexpected values.
5. Seasonality Analysis:  
Use Spark SQL functions or libraries like PySpark's pandas or numpy to analyze seasonality patterns in your data. This could involve time series analysis, trend detection, or Fourier transforms.
6. Anomaly Detection:  
Implement anomaly detection algorithms using Spark SQL or PySpark libraries.
7. Visualization:  
Visualize the results of the analysis using plots, charts, or dashboards to communicate insights effectively.

# About Dataset :

## Climate Change: Earth Surface Temperature Data

Exploring global temperatures since 1750

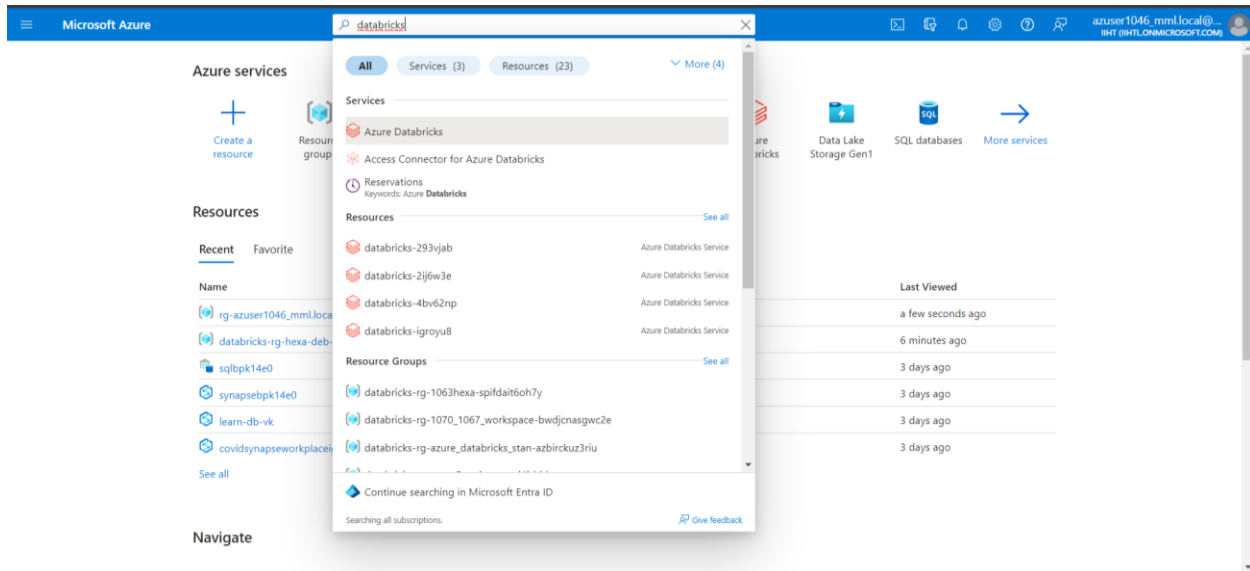
The Berkeley Earth Surface Temperature Study combines 1.6 billion temperature reports from 16 pre-existing archives. It is nicely packaged and allows for slicing into interesting subsets (for example by country). They publish the source data and the code for the transformations they applied. They also use methods that allow weather observations from shorter time series to be included, meaning fewer observations need to be thrown away.

Early data was collected by technicians using mercury thermometers, where any variation in the visit time impacted measurements. In the 1940s, the construction of airports caused many weather stations to be moved. In the 1980s, there was a move to electronic thermometers that are said to have a cooling bias. Our Dataset contains columns Date, Average Temperature, Average Temperature Uncertainty, Country.

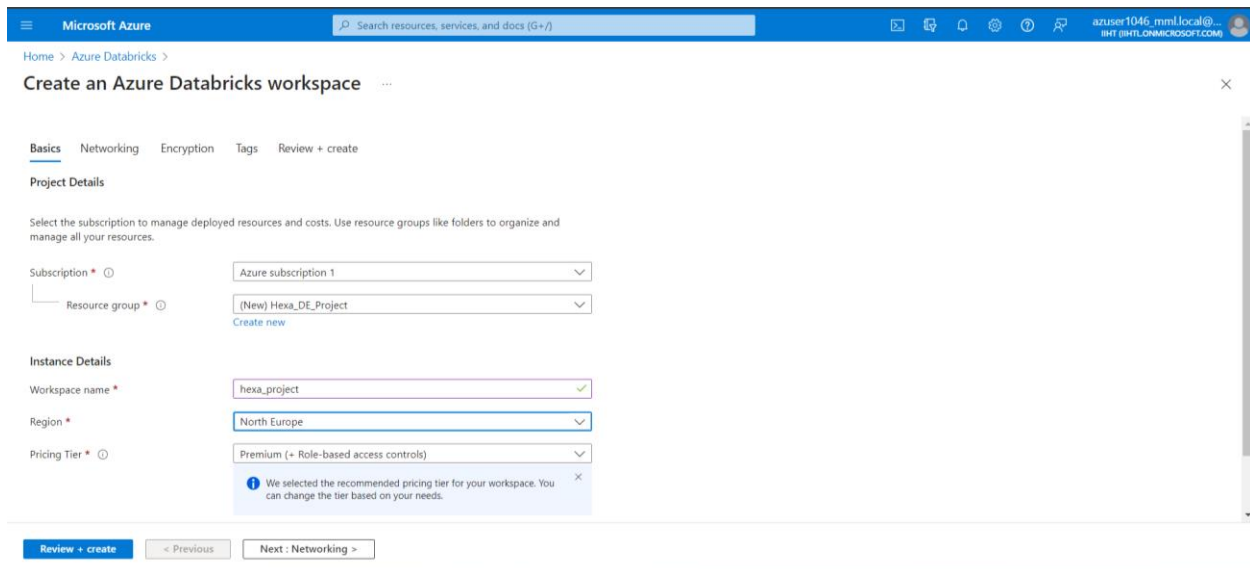
| GlobalLandTemperaturesByCountry.csv  |            |                     |                                       |
|--|------------|---------------------|---------------------------------------|
| D: > Hexaware > Data_Engineering-Batch > Data_Engineering_Projects > Datasets > GlobalLandTemperatures > GlobalLandTemperaturesByCountry.csv |            |                     |                                       |
| 1  | Date       | AverageTemperature  | AverageTemperatureUncertainty,Country |
| 2  | 1743-11-01 | 4.3839999999999995  | 2.294,Åland                           |
| 3  | 1743-12-01 | ,,Åland             |                                       |
| 4  | 1744-01-01 | ,,Åland             |                                       |
| 5  | 1744-02-01 | ,,Åland             |                                       |
| 6  | 1744-03-01 | ,,Åland             |                                       |
| 7  | 1744-04-01 | 1.53,4.68,Åland     |                                       |
| 8  | 1744-05-01 | 6.702000000000001   | 1.789,Åland                           |
| 9  | 1744-06-01 | 11.609000000000002  | 1.577,Åland                           |
| 10   | 1744-07-01 | 15.342,1.41,Åland   |                                       |
| 11   | 1744-08-01 | ,,Åland             |                                       |
| 12   | 1744-09-01 | 11.702,1.517,Åland  |                                       |
| 13   | 1744-10-01 | 5.477,1.862,Åland   |                                       |
| 14   | 1744-11-01 | 3.407,1.425,Åland   |                                       |
| 15   | 1744-12-01 | -2.181,1.641,Åland  |                                       |
| 16   | 1745-01-01 | -3.85,1.841,Åland   |                                       |
| 17   | 1745-02-01 | -6.574999999999998  | 1.36,Åland                            |
| 18   | 1745-03-01 | -4.195,1.213,Åland  |                                       |
| 19   | 1745-04-01 | -0.9660000000000002 | 1.172,Åland                           |
| 20   | 1745-05-01 | ,,Åland             |                                       |
| 21   | 1745-06-01 | ,,Åland             |                                       |
| 22   | 1745-07-01 | ,,Åland             |                                       |
| 23   | 1745-08-01 | ,,Åland             |                                       |
| 24   | 1745-09-01 | ,,Åland             |                                       |
| 25   | 1745-10-01 | ,,Åland             |                                       |
| 26   | 1745-11-01 | ,,Åland             |                                       |
| 27   | 1745-12-01 | ,,Åland             |                                       |
| 28   | 1746-01-01 | ,,Åland             |                                       |
| 29   | 1746-02-01 | ,,Åland             |                                       |
| 30   | 1746-03-01 | ,,Åland             |                                       |
| 31   | 1746-04-01 | ,,Åland             |                                       |
| 32   | 1746-05-01 | ,,Åland             |                                       |
| 33   | 1746-06-01 | ,,Åland             |                                       |
| 34   | 1746-07-01 | ,,Åland             |                                       |
| 35   | 1746-08-01 | ,,Åland             |                                       |
| 36   | 1746-09-01 | ,,Åland             |                                       |
| 37   | 1746-10-01 | ,,Åland             |                                       |

# Tasks Performed :

## Search for Azure Databricks in the resources :



## Created Azure Databricks Workspace :



Microsoft Azure

Search resources, services, and docs (G+/I)

azuser1046\_mml.local@...  
BMT (BMT.LON.MICROSOFT.COM)

Home > Azure Databricks >

Create an Azure Databricks workspace ...

Validation Succeeded

Basics

Networking

Encryption

Tags

Review + create

Summary

Basics

Workspace name  
hexa\_project  
Subscription  
Azure subscription 1  
Resource group  
Hexa\_DE\_Project  
Region  
North Europe  
Pricing Tier  
premium  
Managed Resource Group name

Networking

Deploy Azure Databricks workspace with  
Secure Cluster Connectivity (No Public IP)  
Deploy Azure Databricks workspace in  
your own Virtual Network (VNet)

No  
No

Create

< Previous

Download a template for automation

Microsoft Azure

Search resources, services, and docs (G+/I)

azuser1046\_mml.local@...  
BMT (BMT.LON.MICROSOFT.COM)

Home >

Hexa\_DE\_Project\_hexa\_project | Overview ...

Deployment

Search

Delete Cancel Redeploy Download Refresh

Overview

Inputs

Outputs

Template

Your deployment is complete

Deployment name : Hexa\_DE\_Project\_hexa\_project  
Subscription : Azure subscription 1  
Resource group : Hexa\_DE\_Project

Start time : 25/2/2024, 2:44:12 pm  
Correlation ID : 3a132e70-181e-435a-8b8c-849a0f5f128d

Deployment details

Next steps

Go to resource

Give feedback

Tell us about your experience with deployment

Cost management

Get notified to stay within your budget and prevent unexpected charges on your bill.  
Set up cost alerts >

Microsoft Defender for Cloud

Secure your apps and infrastructure  
Go to Microsoft Defender for Cloud >

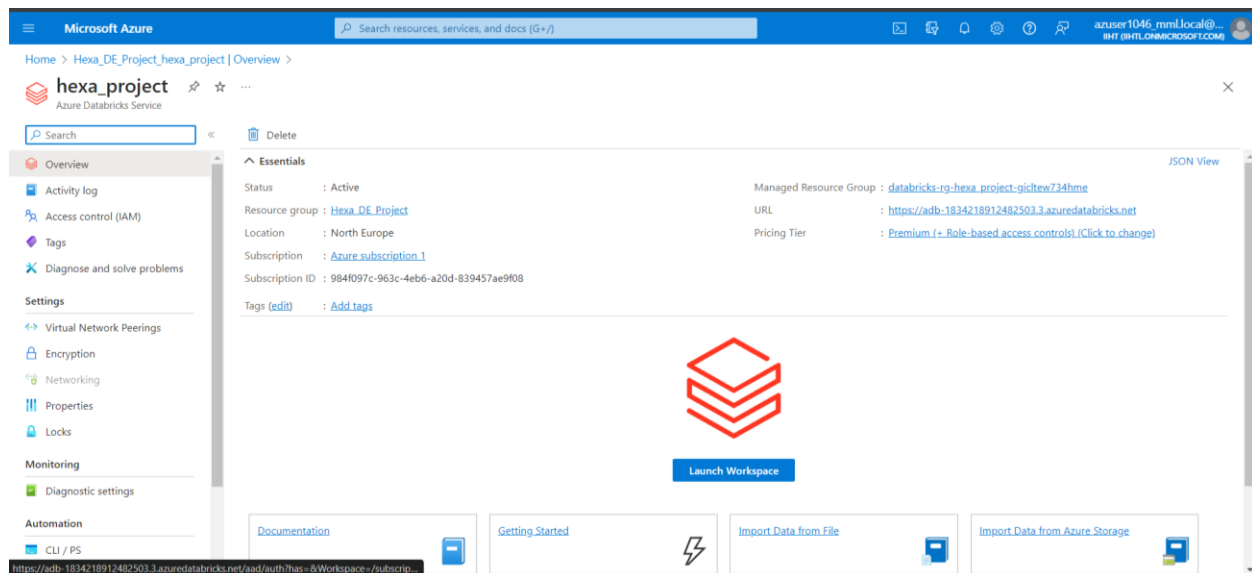
Free Microsoft tutorials

Start learning today >

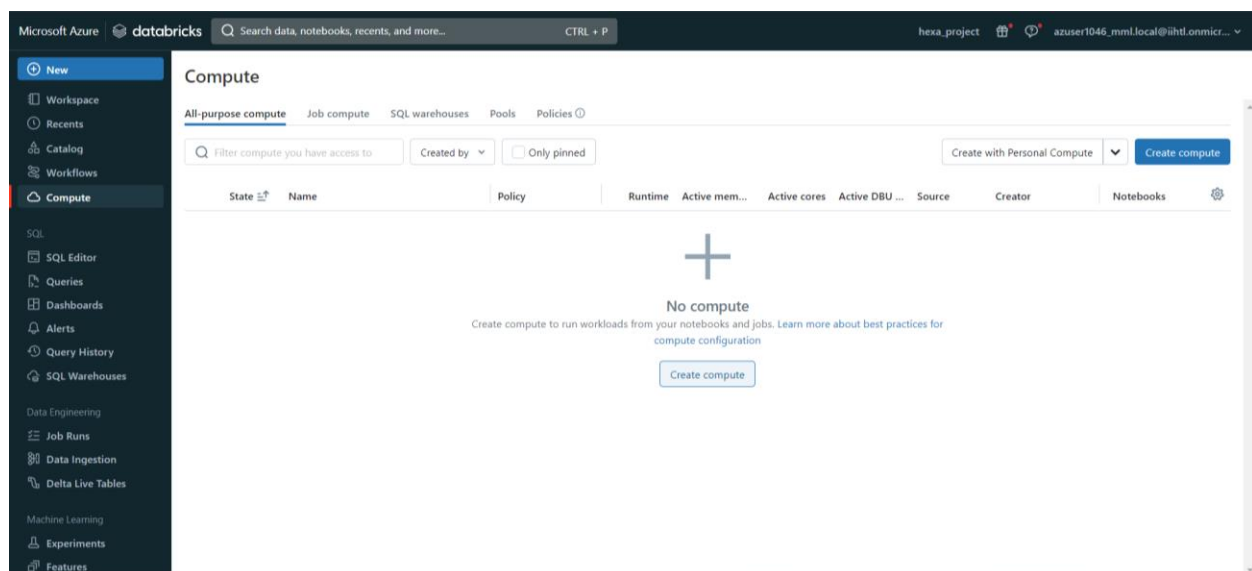
Work with an expert

Azure experts are service provider partners who can help manage your assets on Azure and be your first line of support.  
Find an Azure expert >

## Launch Databricks Workspace :



## Created Cluster in Azure Databricks workspace :





## Configuring the Cluster :

Microsoft Azure databricks Search data, notebooks, recents, and more... CTRL + P hexa\_project azuser1046\_mml.local@iihtl.onmicr...

Compute > Preview Send feedback

azuser1046\_mml.local@iihtl.onmicrosoft.com's Cluster More Terminate Edit

Configuration Notebooks (0) Libraries Event log Spark UI Driver logs Metrics Apps Spark compute UI - Master

Policy Personal Compute

Access mode Single user access Single user azuser1046\_mml.local@iihtl.onmicrosoft...

Performance Databricks Runtime Version 14.3 LTS ML (includes Apache Spark 3.5.0, Scala 2.12)

Use Photon Acceleration

Node type Standard\_DS3\_v2 14 GB Memory, 4 Cores

Terminate after 60 minutes of inactivity

Tags No custom tags

Summary

1 Driver 14 GB Memory, 4 Cores

Runtime 14.3.x-cpu-mi-scala2.12

Standard\_DS3\_v2 0.75 DBU/h

## Created Notebook in Azure Databricks workspace :

Microsoft Azure databricks Notebook s, recents, and more... CTRL + P hexa\_project azuser1046\_mml.local@iihtl.onmicr...

New

Workspace

Recents

Catalog

Workflows

Compute

SQL

SQL Editor

Queries

Dashboards

Alerts

Query History

SQL Warehouses

Data Engineering

Job Runs

Data Ingestion

Delta Live Tables

Machine Learning

Experiments

Features

Repo

Data

File upload

Add data

Compute

Cluster

SQL Warehouse

SQL

Query

Dashboard

Lakeview Dashboard

Alert

Data Engineering

Job

DLT Pipeline

Machine Learning

AutoML Experiment

Experiment

Model

Summary

1 Driver 14 GB Memory, 4 Cores

Runtime 14.3.x-cpu-mi-scala2.12

Standard\_DS3\_v2 0.75 DBU/h

https://adb-1834218912482503.3.azuredatabricks.net/?o=1834218912482503#notebook/cre...

# Created Azure Data Lake Storage account

Microsoft Azure

Search resources, services, and docs (G+)

Home > Storage accounts >

## Create a storage account

Basics | Advanced | Networking | Data protection | Encryption | Tags | Review

**Project details**

Select the subscription in which to create the new storage account. Choose a new or existing resource group to organize and manage your storage account together with other resources.

Subscription \*

Resource group \*   
[Create new](#)

**Instance details**

Storage account name \*

Region \*   
[Deploy to an edge zone](#)

[Review](#) [Previous](#) [Next: Advanced >](#) [Give feedback](#)

portal.azure.com/#view/HubsExtension/DeploymentDetailsBlade/~/overview/id/%2Fsubscriptions%2F984f097c-963c-4eb6-a20d-839457ae9f08%2Fres...

Microsoft Azure

Home >

### adls1046project\_1709099457934 | Overview

Deployment

Search << Delete Cancel Redeploy Download Refresh

**Overview**

**Inputs**

**Outputs**

**Template**

**Your deployment is complete**

Deployment name: adls1046project\_1709099457934  
Subscription: Azure subscription 1  
Resource group: Hexa\_DE\_Project

Start time: 28/2/2024, 11:21:08 am  
Correlation ID: 3695c3e6-b223-479a-b6cb-750c9403b348

**Deployment details**

**Next steps**

[Go to resource](#)

**Give feedback**

[Tell us about your experience with deployment](#)

**Cost Management**

Get notified to stay within your budget and prevent unexpected charges on your bill.  
[Set up cost alerts >](#)

**Microsoft Defender for Cloud**

Secure your apps and infrastructure  
[Go to Microsoft Defender for Cloud >](#)

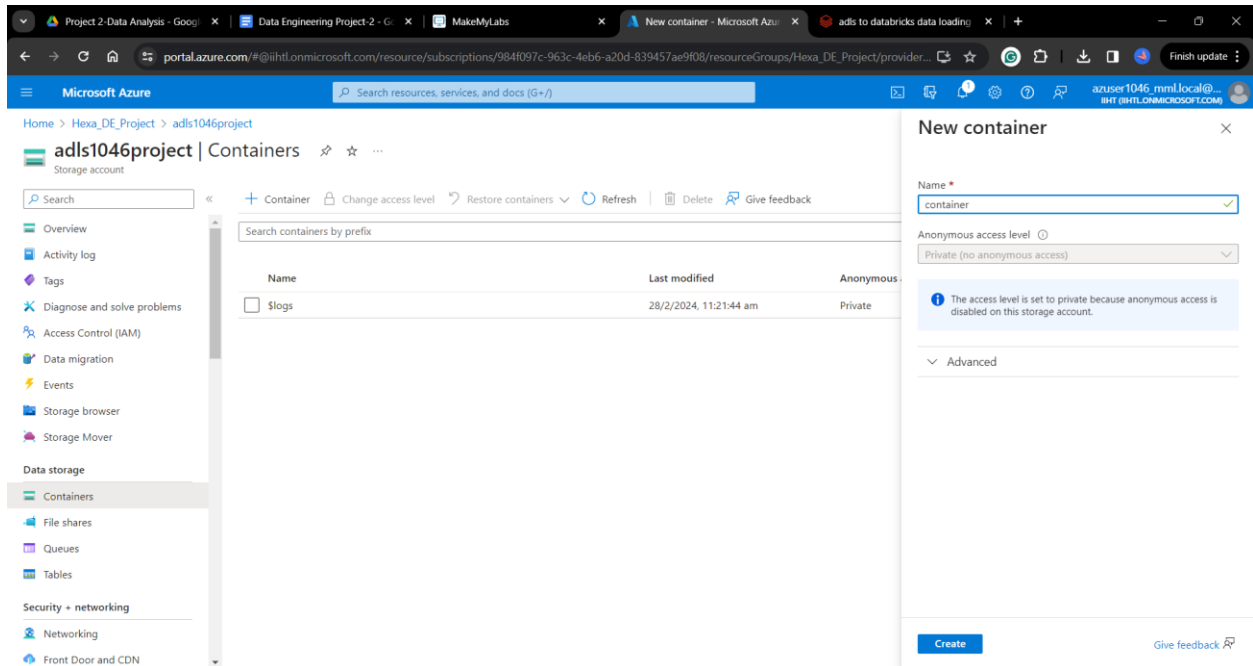
**Free Microsoft tutorials**

[Start learning today >](#)

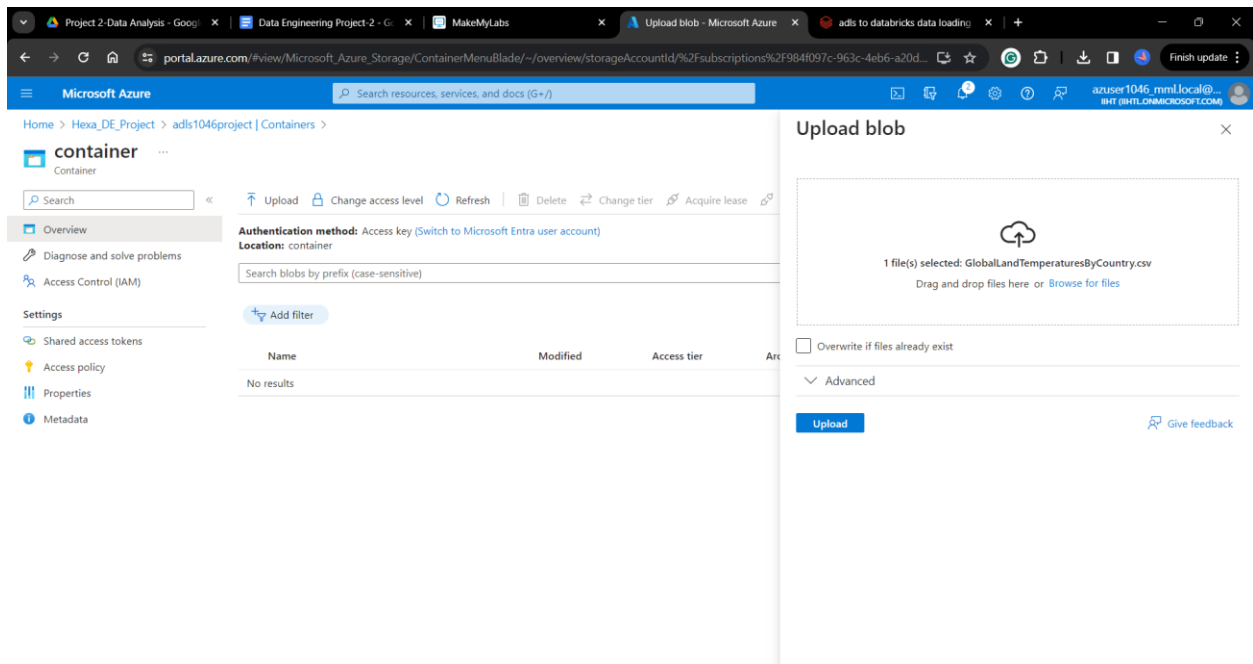
**Work with an expert**

Azure experts are service provider partners who can help manage your assets on Azure and be your first line of support.  
[Find an Azure expert >](#)

## Created container in ADLS



## Uploading Dataset in the ADLS



## Data Preparation:

- Loaded the global temperature data into Spark DataFrame from our data source Azure Blob Storage.

### Load Data from Data Sources:

- For Azure Blob Storage: Use the **`spark.read.format("csv").load("wasbs://<container-name>@<storage-account-name>.blob.core.windows.net/<path>")`** method to load CSV files from Azure Blob Storage.
- Once the data is loaded and formatted, created a Spark DataFrame to represent the data in a structured format for further analysis.
- Used the **`spark.createDataFrame()`** method or DataFrame APIs to create the DataFrame from the loaded data.

```
Just now (19s) Cell 1 Python
```

```
storage_account_name = "adls1046project"
container_name = "container"
file_name = "GlobalLandTemperaturesByCountry.csv"
storage_account_access_key = "KUKAhIfCYt/M/1Co79hLgWIsSzICBxQ7TJJ3s7DwtrO1gZWe/lRA0Rh43+Qmu1+PR31WfvqtRx6w+ASthC50fg=="
spark.conf.set("fs.azure.account.key." + storage_account_name + ".blob.core.windows.net", storage_account_access_key)
df = spark.read.format("csv").load(f"wasbs://{container_name}@{storage_account_name}.blob.core.windows.net/{file_name}", inferSchema=True,
header=True)
df.show()
```

▶ (3) Spark Jobs

df: pyspark.sql.dataframe.DataFrame = [Date: date, AverageTemperature: double ... 2 more fields]

| Date       | AverageTemperature | AverageTemperatureUncertainty | Country |
|------------|--------------------|-------------------------------|---------|
| 1743-11-01 | 4.3839999999999995 | 2.294                         | Åland   |
| 1743-12-01 | NULL               | NULL                          | Åland   |
| 1744-01-01 | NULL               | NULL                          | Åland   |
| 1744-02-01 | NULL               | NULL                          | Åland   |
| 1744-03-01 | NULL               | NULL                          | Åland   |
| 1744-04-01 | 1.53               | 4.68                          | Åland   |
| 1744-05-01 | 6.7020000000000001 | 1.789                         | Åland   |
| 1744-06-01 | 11.609000000000002 | 1.577                         | Åland   |
| 1744-07-01 | 15.342             | 1.41                          | Åland   |
| 1744-08-01 | NULL               | NULL                          | Åland   |
| 1744-09-01 | 11.702             | 1.517                         | Åland   |
| 1744-10-01 | 5.477              | 1.862                         | Åland   |
| 1744-11-01 | 3.407              | 1.425                         | Åland   |
| 1744-12-01 | -2.181             | 1.641                         | Åland   |

## Data Exploration:

To explore our data using Spark SQL, We followed these steps:

1. Load the data into Spark DataFrame.
2. Used Spark SQL queries to explore the data, including identifying the structure, schema, and basic statistics.

```
▶ ▾ ✓ Just now (<1s) Cell 2 Python ✨ ⌵ ⋮
```

```
# Identify the structure and schema of the dataset
print("Schema of the dataset:")
data.printSchema()
```

```
Schema of the dataset:
root
|-- Date: date (nullable = true)
|-- AverageTemperature: double (nullable = true)
|-- AverageTemperatureUncertainty: double (nullable = true)
|-- Country: string (nullable = true)
```

```
▶ ▾ Cell 4 Python ✨ ⌵ ⋮
```

```
# Use Spark SQL queries to explore the data
# Example 1: Count the number of records in the dataset
record_count = spark.sql("SELECT COUNT(*) AS record_count FROM climate_data").collect()[0]["record_count"]
print("Number of records in the dataset:", record_count)
```

```
Number of records in the dataset: 577462
```

```
▶ ▾ ✓ Just now (2s) Cell 6 Python ✨ ⌵ ⋮
```

```
# Example 3: Display distinct values in a particular column
print("Distinct Countries:")
spark.sql("SELECT DISTINCT Country FROM climate_data").show()
```

▶ (2) Spark Jobs

```
Distinct Countries:
+-----+
| Country|
+-----+
| Anguilla|
| Åland|
| Afghanistan|
| Africa|
| Algeria|
| Argentina|
| Angola|
| Baker Island|
| Albania|
| Bahamas|
| American Samoa|
| Andorra|
| Antarctica|
| Aruba|
| Azerbaijan|
| Armenia|
| Asia|
```

```

Just now (3s) Cell 7 Python
# Example 4: Calculate the average value of a numerical column
print("Average value of a of temperature of the year:")
avg_temp = spark.sql("SELECT DATE_PART('YEAR', Date), AVG(AverageTemperature) AS avg_temp FROM climate_data GROUP BY DATE_PART('YEAR', Date) ORDER BY DATE_PART('YEAR', Date)")
avg_temp.show()

```

▶ (2) Spark Jobs

avg\_temp: pyspark.sql.dataframe.DataFrame = [date\_part(YEAR, Date): integer, avg\_temp: double]

Average value of a of temperature of the year:

| date_part(YEAR, Date) | avg_temp           |
|-----------------------|--------------------|
| 1743                  | 5.18414            |
| 1744                  | 9.8378975          |
| 1745                  | 1.3871250000000004 |
| 1746                  | NULL               |
| 1747                  | NULL               |
| 1748                  | NULL               |
| 1749                  | NULL               |
| 1750                  | 9.129352727272728  |
| 1751                  | 9.167387499999998  |
| 1752                  | 4.413386666666668  |
| 1753                  | 8.870820754716977  |
| 1754                  | 8.822018957345971  |
| 1755                  | 8.530536277602524  |
| 1756                  | 9.17988625592417   |
| 1757                  | 8.993332283464566  |
| 1758                  | 8.13037054263566   |
| 1759                  | 0.2612577160000001 |

## Data Cleaning:

- Used Spark SQL queries or DataFrame operations to perform data cleaning operations to inspect the data and to handle missing values, and any inconsistencies in the data.

```

Cell 2
# Handle missing values
# Drop rows with any missing temperature values
cleaned_data = data.dropna(subset=["AverageTemperature"])
cleaned_data.show()

```

| Date       | AverageTemperature | AverageTemperatureUncertainty | Country |
|------------|--------------------|-------------------------------|---------|
| 1743-11-01 | 4.384              | 2.294                         | Aland   |
| 1744-04-01 | 1.53               | 4.68                          | Aland   |
| 1744-05-01 | 6.702              | 1.789                         | Aland   |
| 1744-06-01 | 11.609             | 1.577                         | Aland   |
| 1744-07-01 | 15.342             | 1.41                          | Aland   |
| 1744-09-01 | 11.702             | 1.517                         | Aland   |
| 1744-10-01 | 5.477              | 1.862                         | Aland   |
| 1744-11-01 | 3.407              | 1.425                         | Aland   |
| 1744-12-01 | -2.181             | 1.641                         | Aland   |
| 1745-01-01 | -3.85              | 1.841                         | Aland   |
| 1745-02-01 | -6.575             | 1.36                          | Aland   |
| 1745-03-01 | -4.195             | 1.213                         | Aland   |
| 1745-04-01 | -0.966             | 1.172                         | Aland   |
| 1750-01-01 | 1.091              | 1.119                         | Aland   |
| 1750-02-01 | 0.809              | 3.353                         | Aland   |
| 1750-03-01 | 0.923              | 4.716                         | Aland   |
| 1750-04-01 | 3.943              | 1.434                         | Aland   |
| 1750-05-01 | 6.265              | 1.339                         | Aland   |

▶

Cell 3

Python

+

+

⋮

```
# Drop rows with any missing date values
cleaned_data = cleaned_data.dropna(subset=["Date"])
cleaned_data.show()
```

| Date       | AverageTemperature | AverageTemperatureUncertainty | Country |
|------------|--------------------|-------------------------------|---------|
| 1743-11-01 | 4.384              | 2.294                         | Åland   |
| 1744-04-01 | 1.53               | 4.68                          | Åland   |
| 1744-05-01 | 6.702              | 1.789                         | Åland   |
| 1744-06-01 | 11.609             | 1.577                         | Åland   |
| 1744-07-01 | 15.342             | 1.41                          | Åland   |
| 1744-09-01 | 11.702             | 1.517                         | Åland   |
| 1744-10-01 | 5.477              | 1.862                         | Åland   |
| 1744-11-01 | 3.407              | 1.425                         | Åland   |
| 1744-12-01 | -2.181             | 1.641                         | Åland   |
| 1745-01-01 | -3.85              | 1.841                         | Åland   |
| 1745-02-01 | -6.575             | 1.36                          | Åland   |
| 1745-03-01 | -4.195             | 1.213                         | Åland   |
| 1745-04-01 | -0.966             | 1.172                         | Åland   |
| 1750-01-01 | 1.091              | 1.119                         | Åland   |
| 1750-02-01 | 0.809              | 3.353                         | Åland   |
| 1750-03-01 | 0.923              | 4.716                         | Åland   |
| 1750-04-01 | 3.943              | 1.434                         | Åland   |
| 1750-05-01 | 6.265              | 1.339                         | Åland   |

▶

Just now (1s)

Cell 4

Python

+

+

⋮

```
# Remove rows with missing country values
from pyspark.sql.functions import col
cleaned_data = cleaned_data.filter(col("Country").isNotNull())
cleaned_data.show()
```

▶ (1) Spark Jobs

▶ cleaned\_data: pyspark.sql.dataframe.DataFrame = [Date: date, AverageTemperature: double ... 2 more fields]

| Date       | AverageTemperature  | AverageTemperatureUncertainty | Country |
|------------|---------------------|-------------------------------|---------|
| 1743-11-01 | 4.3839999999999995  | 2.294                         | Åland   |
| 1744-04-01 | 1.53                | 4.68                          | Åland   |
| 1744-05-01 | 6.702000000000001   | 1.789                         | Åland   |
| 1744-06-01 | 11.609000000000002  | 1.577                         | Åland   |
| 1744-07-01 | 15.342              | 1.41                          | Åland   |
| 1744-09-01 | 11.702              | 1.517                         | Åland   |
| 1744-10-01 | 5.477               | 1.862                         | Åland   |
| 1744-11-01 | 3.407               | 1.425                         | Åland   |
| 1744-12-01 | -2.181              | 1.641                         | Åland   |
| 1745-01-01 | -3.85               | 1.841                         | Åland   |
| 1745-02-01 | -6.574999999999998  | 1.36                          | Åland   |
| 1745-03-01 | -4.195              | 1.213                         | Åland   |
| 1745-04-01 | -0.9660000000000002 | 1.172                         | Åland   |
| 1750-01-01 | 1.0910000000000006  | 1.119                         | Åland   |
| 1750-02-01 | 0.8090000000000002  | 3.353                         | Åland   |
| 1750-03-01 | 0.9229999999999999  | 4.716                         | Åland   |
| 1750-04-01 | 3.943               | 1.4340000000000002            | Åland   |
| 1750-05-01 | 6.2650000000000015  | 1.339                         | Åland   |

```

Cell 5
Python
# Drop irrelevant columns
cleaned_data = cleaned_data.drop("AverageTemperatureUncertainty")
cleaned_data.show()

```

|            |        |       |
|------------|--------|-------|
| 1744-05-01 | 6.702  | Åland |
| 1744-06-01 | 11.609 | Åland |
| 1744-07-01 | 15.342 | Åland |
| 1744-09-01 | 11.702 | Åland |
| 1744-10-01 | 5.477  | Åland |
| 1744-11-01 | 3.407  | Åland |
| 1744-12-01 | -2.181 | Åland |
| 1745-01-01 | -3.85  | Åland |
| 1745-02-01 | -6.575 | Åland |
| 1745-03-01 | -4.195 | Åland |
| 1745-04-01 | -0.966 | Åland |
| 1750-01-01 | 1.091  | Åland |
| 1750-02-01 | 0.809  | Åland |
| 1750-03-01 | 0.923  | Åland |
| 1750-04-01 | 3.943  | Åland |
| 1750-05-01 | 6.265  | Åland |
| 1750-06-01 | 12.408 | Åland |
| 1750-07-01 | 16.683 | Åland |

only showing top 20 rows

## Error Detection:

- Utilized Spark SQL queries or built-in functions to detect errors in your data. This involved identifying inconsistencies, duplicates, or unexpected values.

```

# Check for missing values in relevant columns (e.g., Country, AverageTemperature)
missing_values = data.filter((col("Country").isNull()) | (col("AverageTemperature").isNull()))
if missing_values.count() > 0:
    print("Missing values found in the dataset:")
    missing_values.show()
else:
    print("No missing values found in the dataset.")

```

Missing values found in the dataset:

| Date       | AverageTemperature | AverageTemperatureUncertainty | Country |
|------------|--------------------|-------------------------------|---------|
| 1743-12-01 | null               | null                          | Åland   |
| 1744-01-01 | null               | null                          | Åland   |
| 1744-02-01 | null               | null                          | Åland   |
| 1744-03-01 | null               | null                          | Åland   |
| 1744-08-01 | null               | null                          | Åland   |
| 1745-05-01 | null               | null                          | Åland   |
| 1745-06-01 | null               | null                          | Åland   |
| 1745-07-01 | null               | null                          | Åland   |
| 1745-08-01 | null               | null                          | Åland   |
| 1745-09-01 | null               | null                          | Åland   |
| 1745-10-01 | null               | null                          | Åland   |
| 1745-11-01 | null               | null                          | Åland   |
| 1745-12-01 | null               | null                          | Åland   |
| 1746-01-01 | null               | null                          | Åland   |
| 1746-02-01 | null               | null                          | Åland   |
| 1746-03-01 | null               | null                          | Åland   |
| 1746-04-01 | null               | null                          | Åland   |



Cell 3

```
# Check for duplicates
duplicate_count = data.groupBy(data.columns).count().filter("count > 1").count()
if duplicate_count > 0:
    print("Duplicates found in the dataset.")
else:
    print("No duplicates found in the dataset.")
```

No duplicates found in the dataset.

```
# Check for outliers or unrealistic values in AverageTemperature
# you can set a threshold to identify outliers
min_threshold = -30
max_threshold = 50
outliers = data.filter((col("AverageTemperature") < min_threshold) | (col("AverageTemperature") > max_threshold))
if outliers.count() > 0:
    print("Outliers found in the AverageTemperature column:")
    outliers.show()
else:
    print("No outliers found in the AverageTemperature column.")
```

Outliers found in the AverageTemperature column:

| Date       | AverageTemperature | AverageTemperatureUncertainty | Country |
|------------|--------------------|-------------------------------|---------|
| 1823-02-01 | -31.746            | 3.438                         | Denmark |
| 1831-02-01 | -31.28             | 4.364                         | Denmark |
| 1832-02-01 | -31.259            | 3.57                          | Denmark |
| 1834-01-01 | -31.572            | 3.27                          | Denmark |
| 1835-02-01 | -31.831            | 3.245                         | Denmark |
| 1836-02-01 | -32.625            | 3.677                         | Denmark |
| 1838-01-01 | -31.057            | 3.493                         | Denmark |
| 1839-01-01 | -31.251            | 3.188                         | Denmark |
| 1839-12-01 | -31.045            | 2.886                         | Denmark |
| 1840-01-01 | -31.082            | 3.549                         | Denmark |
| 1841-01-01 | -31.097            | 2.931                         | Denmark |
| 1844-02-01 | -31.911            | 3.263                         | Denmark |
| 1845-02-01 | -31.081            | 3.222                         | Denmark |
| 1845-12-01 | -31.152            | 3.11                          | Denmark |
| 1848-02-01 | -31.245            | 3.602                         | Denmark |

## Seasonality Analysis:

- Use Spark SQL functions or libraries like PySpark's pandas or numpy to analyze seasonality patterns in your data.

```
Cell 2 Python
```

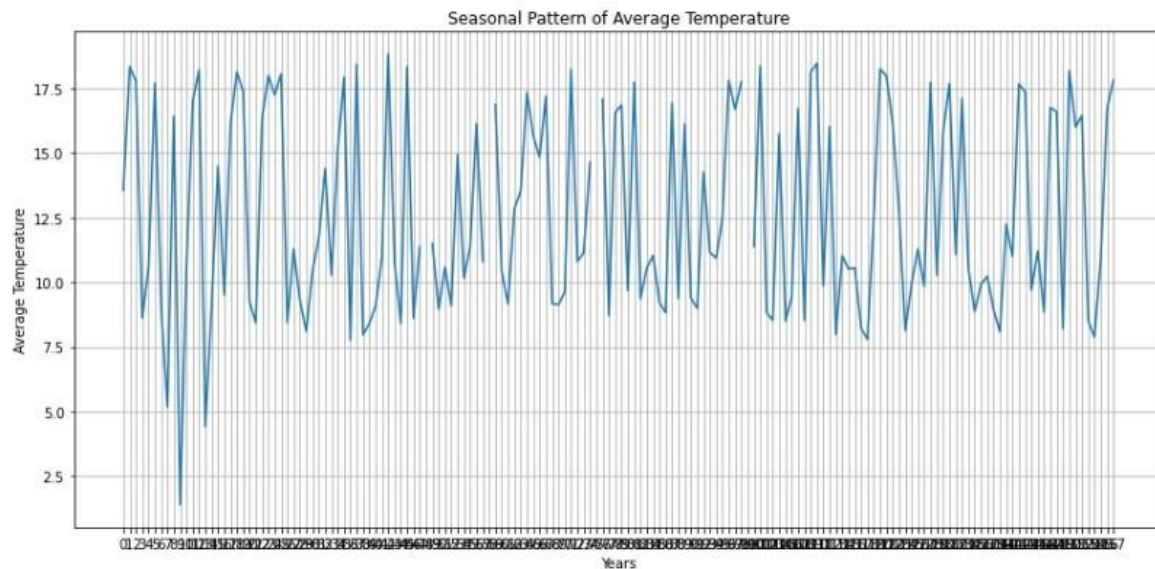
```
# Ensure 'Date' column is of type DateType
data = data.withColumn("Date", data["Date"].cast("date"))
data.show()
```

| Date       | AverageTemperature | AverageTemperatureUncertainty | Country |
|------------|--------------------|-------------------------------|---------|
| 1743-11-01 | 4.384              | 2.294                         | Aland   |
| 1743-12-01 | null               | null                          | Aland   |
| 1744-01-01 | null               | null                          | Aland   |
| 1744-02-01 | null               | null                          | Aland   |
| 1744-03-01 | null               | null                          | Aland   |
| 1744-04-01 | 1.53               | 4.68                          | Aland   |
| 1744-05-01 | 6.702              | 1.789                         | Aland   |
| 1744-06-01 | 11.609             | 1.577                         | Aland   |
| 1744-07-01 | 15.342             | 1.41                          | Aland   |
| 1744-08-01 | null               | null                          | Aland   |
| 1744-09-01 | 11.702             | 1.517                         | Aland   |
| 1744-10-01 | 5.477              | 1.862                         | Aland   |
| 1744-11-01 | 3.407              | 1.425                         | Aland   |
| 1744-12-01 | -2.181             | 1.641                         | Aland   |
| 1745-01-01 | -3.85              | 1.841                         | Aland   |
| 1745-02-01 | -6.575             | 1.36                          | Aland   |
| 1745-03-01 | -4.195             | 1.213                         | Aland   |
| 1745-04-01 | -0.966             | 1.172                         | Aland   |

```
Cell 6
```

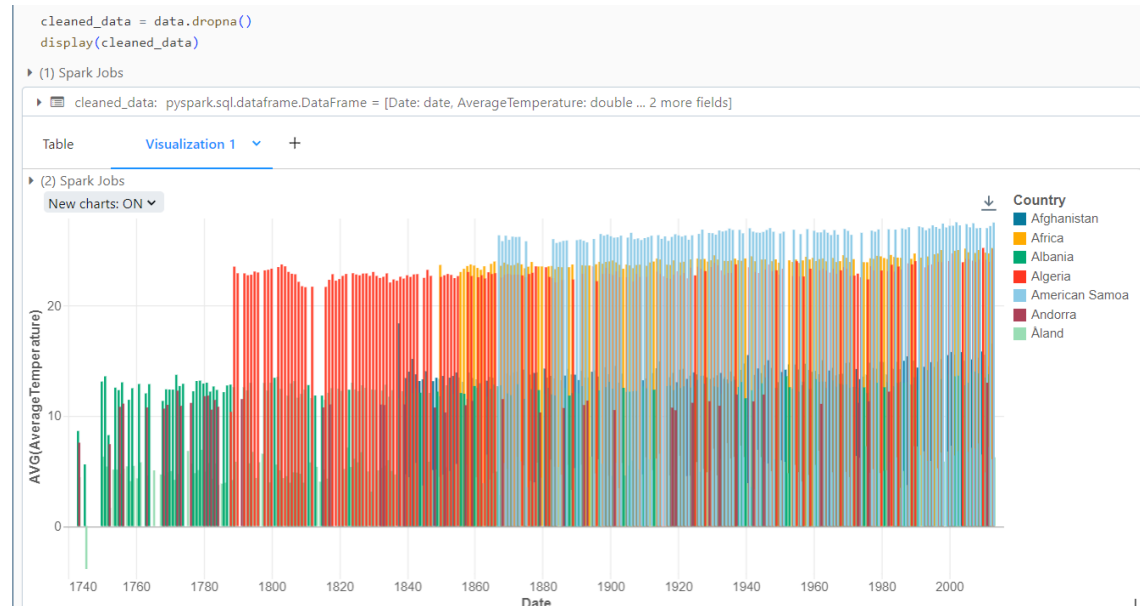
```
# Convert Spark DataFrame to Pandas DataFrame for visualization
yearly_avg_temp_pandas = yearly_avg_temp.toPandas()

# Plot the seasonal pattern of average temperature
plt.figure(figsize=(12, 6))
plt.plot(yearly_avg_temp_pandas["AvgTemperature"])
plt.title("Seasonal Pattern of Average Temperature")
plt.xlabel("Years")
plt.ylabel("Average Temperature")
plt.xticks(range(len(yearly_avg_temp_pandas)))
plt.grid(True)
plt.tight_layout()
plt.show()
```



## Visualization:

- Visualized analysis results using Databricks' built-in visualization tools.



## Data Analysis:

- Utilized PySparkSQL queries to analyze the data for errors, seasonality, and anomalies.
- Aggregate the data to calculate statistics such as average.

```
3 days ago (1s) Cell 3 Python
```

```
print("Total Countries:")
spark.sql("SELECT Count(DISTINCT Country) as Total_Countries FROM climate_data").show()
```

(3) Spark Jobs

```
Distinct Countries:
+-----+
|Total_Countries|
+-----+
|          243|
+-----+
```

3 days ago (1s) Cell 4

```
avg_temp = spark.sql("SELECT DATE_PART('YEAR', dt) as Year, AVG(AverageTemperature) AS avg_temp_by_year, Country FROM climate_data GROUP BY DATE_PART('YEAR', dt), Country ORDER BY Country, Year")
display(avg_temp.dropna())
```


(2) Spark Jobs

avg\_temp: pyspark.sql.dataframe.DataFrame = [Year: integer, avg\_temp\_by\_year: double ... 1 more field]

Table Avg\_Temp\_by\_Year + New result table: OFF

|   | Year | avg_temp_by_year   | Country     |
|---|------|--------------------|-------------|
| 1 | 1838 | 18.379571428571428 | Afghanistan |
| 2 | 1840 | 13.413454545454545 | Afghanistan |
| 3 | 1841 | 13.9976            | Afghanistan |
| 4 | 1842 | 15.154666666666667 | Afghanistan |
| 5 | 1843 | 13.75625           | Afghanistan |
| 6 | 1844 | 13.148750000000001 | Afghanistan |
| 7 | 1845 | 13.305833333333332 | Afghanistan |

10,000 rows | Truncated data | 1.19 seconds runtime Refreshed 3 days ago



Average temperature by year

3 days ago (<1s) Cell 5

```
avg_temp.dropna().createOrReplaceTempView("Avg_temp")
```

3 days ago (1s) Cell 6 Python

```
avg_data = spark.sql("SELECT Avg(avg_temp_by_year) as avg_temp_by_Country, Country from Avg_temp GROUP BY Country ORDER BY Country")
display(avg_data)
```

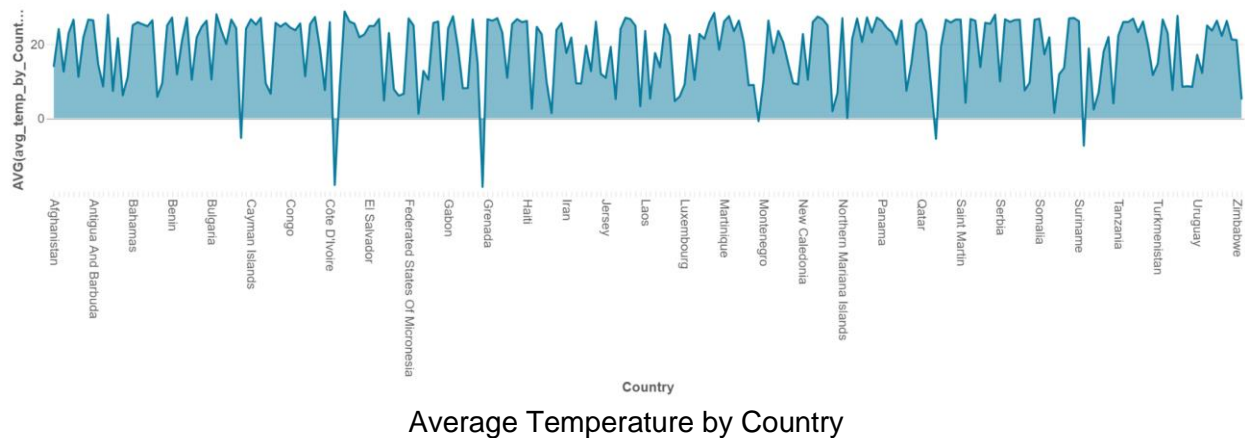
(5) Spark Jobs

avg\_data: pyspark.sql.dataframe.DataFrame = [avg\_temp\_by\_Country: double, Country: string]

Table Avg\_temp\_by\_country + New result table: OFF

|   | avg_temp_by_Country | Country        |
|---|---------------------|----------------|
| 1 | 14.061307148423012  | Afghanistan    |
| 2 | 24.07212044978083   | Africa         |
| 3 | 12.577367012824885  | Albania        |
| 4 | 22.939589924576175  | Algeria        |
| 5 | 26.602235920614174  | American Samoa |
| 6 | 11.171873552945188  | Andorra        |
| 7 | 21.818182360779808  | Anaola         |

242 rows | 1.22 seconds runtime Refreshed 3 days ago



## Conclusion:

In conclusion, the project successfully achieved its objectives of analyzing the Global Land Average Temperature dataset using PySpark on Azure Databricks. By leveraging distributed computing capabilities and advanced analytics techniques, the project provided actionable insights into temperature trends and anomalies, contributing to efforts to understand and address climate change.

The project demonstrates the power of using cloud-based data analytics platforms like Azure Databricks and PySpark for processing and analyzing large-scale environmental datasets, enabling data-driven decision-making and fostering scientific research in climate science and related fields.

## References:

<https://learn.microsoft.com/en-us/azure/storage/blobs/create-data-lake-storage-account>  
<https://learn.microsoft.com/en-in/azure/synapse-analytics/get-started-create-workspace>  
<https://blog.arinti.be/databricks-importing-data-from-a-blob-storage-2b8dc700d029>  
<https://docs.databricks.com/en/ extras/notebooks/source/data-import/azure-blob-store.html>